

# Unified-OneHead Multi-Task Challenge

Object Detection + Semantic Segmentation + Image Classification with a **single head**

---

## o. TL;DR

Key Point	Requirement
<b>Task</b>	Build <b>one head (2-3 layers)</b> that outputs detection + segmentation + classification simultaneously
<b>Data</b>	Three <i>mutually exclusive</i> mini-datasets (< 120 MB total) supplied by the course; one-command download
<b>Hardware</b>	Free Google Colab GPU (T4 or V100) • <b>Total training time <math>\leq</math> 2 h</b>
<b>Core difficulty</b>	<i>Catastrophic forgetting</i> : after alternating training, each task must drop $\leq 5\%$ vs. its single-task baseline
<b>Deliverables</b>	colab.ipynb • report.md • llm_dialogs.zip (all LLM chat logs)
<b>Grading</b>	20 + 20 + 25(+5) + 10 + 15 + 10 = 100 pts
<b>Deadline</b>	<b>XX/YY (Fri) 23:59 GMT+8</b> — submit via GitHub pull request

---

## 1. Official Mini-Datasets

Task	Subset Name	Origin	Download Size	Images (train/val)	Annotation
Detection	<b>Mini-COCO-Det</b>	COCO 2017 (10 classes)	45 MB	300 (240 / 60)	COCO JSON
Segmentation	<b>Mini-VOC-Seg</b>	PASCAL VOC 2012	30 MB	300 (240 / 60)	PNG masks
Classification	<b>Imagenette-160</b>	Imagenette v2	25 MB	300 (240 / 60)	folder / label

All image files are **non-overlapping across the three subsets. Hidden test sets (60 images per task) will be used by the TAs for final evaluation.**

Dataset directory layout after extraction:

```
data/
├─ mini_coco_det/{train,val}
├─ mini_voc_seg/{train,val}
└─ imagenette_160/{train,val}
```

## 2. Model Specifications

Component	Constraint
<b>Backbone</b>	Choose <b>one</b> : MobileNetV3-Small (2.5 M), EfficientNet-Bo (5.3 M), YOLOv8-n <i>backbone</i> (3.2 M), or Fast-SCNN (1.1 M). Public ImageNet weights may be loaded.
<b>Neck</b>	$\leq 2$ conv/BN/ReLU layers <i>or</i> a single FPN layer.
<b>Head</b>	Exactly <b>2–3 layers</b> (Conv or MLP) and <b>single branch</b> : must emit <i>all three</i> outputs at once. Example output schema (stride 16 feature map): <ul style="list-style-type: none"><li>• Detection: <math>N \times (cx, cy, w, h, conf, C_{det})</math></li><li>• Segmentation: <math>C_{seg} \times H \times W</math> mask</li><li>• Classification: <math>C_{cls}</math> logits</li></ul>
<b>Total parameters</b>	$< 8 \text{ M}$ (entire model).
<b>Inference speed</b>	$\leq 150 \text{ ms}$ per $512 \times 512$ image on Colab T4.

## 3. Training Schedule & Forgetting Criterion

Stage 0 (optional) – warm-up / ImageNet pretrain  
Stage 1 – Train ONLY on Dataset A (Seg) → record mIoU\_base  
Stage 2 – Train ONLY on Dataset B (Det) → measure mIoU\_drop  
Stage 3 – Train ONLY on Dataset C (Cls) → measure mIoU\_drop + mAP\_drop

Required after Stage 3:

mIoU  $\geq$  mIoU\_base - 5 %  
mAP  $\geq$  mAP\_base - 5 %  
Top-1  $\geq$  Top1\_base - 5 %

### 3.1 Forgetting-mitigation toolbox (pick any)

- Elastic Weight Consolidation (EWC)
- Learning without Forgetting (LwF)
- Replay buffer ( $\leq 10$  images per task per step)
- Knowledge distillation (previous checkpoint as teacher)

## 4. Scoring (100 pts)

Item	Description	Pts
<b>Design &amp; motivation</b>	Soundness and creativity of single-head architecture	20
<b>Training schedule &amp; forgetting remedy</b>	Completeness and theoretical justification	20
<b>Performance</b>	All tasks within the 5 % drop $\rightarrow$ 25 pts <i>Bonus + 5 pts</i> if <b>every</b> metric $\geq$ its baseline	25 (+5)
<b>Resource efficiency</b>	Training $\leq 2$ h, params $< 8$ M, inference $< 150$ ms	10
<b>Report quality</b>	Clarity, figures, citations	15
<b>LLM dialogue logs</b>	Complete prompts + responses	10

## 5. Deliverables

```
your_repo/
├── colab.ipynb      ← runnable end-to-end notebook
├── report.md        ← architecture, schedule, results, analysis
├── llm_dialogs.zip  ← ALL ChatGPT / LLM logs (txt or json)
├── README.md        ← reproduction steps & requirements
└── scripts/         ← helper & evaluation scripts
```

*Submit a GitHub pull request to the course repo including the commit hash.*

## 6. Evaluation Command (used by TAs)

```
python eval.py \  
  --weights your_model.pt \  
  --data_root data \  
  --tasks all          # outputs mAP, mIoU, Top-1
```

---

## 7. FAQ

Question	Answer
<b>Can I use a different backbone?</b>	Yes, as long as the <i>total</i> parameter count is < 8 M and you justify your choice in the report.
<b>May I implement the head as a small Transformer?</b>	Yes— $\leq 3$ layers, still within the parameter limit.
<b>How do I balance the three losses?</b>	Up to you; describe your search strategy in the report.
<b>Are LLM logs really mandatory?</b>	Yes. Missing or incomplete logs = 0 pts for that category.
<b>Can I train on multiple GPUs or paid Colab?</b>	No. The TA validation uses a single free T4/V100. Your code must reproduce results in that environment.

---

## 8. Academic Integrity

- **Provide all LLM conversations** in `llm_dialogs.zip`.
- External code/data must be cited.
- Plagiarism or undisclosed extra data leads to point deductions or failure.